

NetFlow to guard the infrastructure

Yann Berthier

NANOG39 - Feb 4-7, 07

Outline I

1 An introductory view of NetFlow

2 The flow cache

- Overview
- Populating the cache
- Keeping the cache under control
- Querying the cache in real time
- Exporting the content of the cache

3 Collect & analysis

- Overview
- Storage
- Tools
- nfdump / NfSen
- direction & biflows
- From PCAP to NetFlow
- NetFlow and BGP
- Other sources

Outline II

- Reductio ab aggregatio

4 What's next

- Perhaps not really next
- Resources

Outline

- 1 An introductory view of NetFlow
- 2 The flow cache
- 3 Collect & analysis
- 4 What's next

What is NetFlow ?

- A technology to gather informations on forwarded packets
 - In one or several caches
 - Whose content can be queried real time for troubleshooting
 - And exported to *collectors*
- A protocol to export those records to collectors
 - To use sophisticated tools for analysis
 - Near real-time analysis
 - For long-term analysis, trending, whatever
 - To store and retrieve if needed (forensics) / when asked
- Comes from one vendor, but industry + IETF standard (not quite, but ...)
 - Other *concurrent* technologies: CRANE, Diameter, LFAP, IPDR, sFlow, ...

What is a flow ?

A loose definition

A set of packets having common characteristics

Definition

A flow is a unidirectional set of packets that arrive at the router on the same subinterface, have the same source and destination IP addresses, Layer 4 protocol, TCP/UDP source and destination ports, and the same ToS (type of service) byte in the IP headers

Anatomy of a flow (v5)

- 7-tuple *key* fields
 - saddr, daddr, sport, dport, L3 proto, ToS, input iflIndex
- Additional fields
 - Byte count, packet count, start time, end time, output iflIndex, TCP flags, next hop, src AS, dst AS

Some characteristics

- L3-L4
- Flows are *unidirectionals* (ingress / egress)
- Flow-cache comes before ACL lookup
- Comes at a cost (memory, CPU) for the router
- Not the perfect solution - but not a lot of other candidates either

Rationale

- Old techno (96)
- A lot happening in the protocol and implementation front
- A lot of interest and projects evolving around NetFlow
- Questions keep being raised (How do i monitor ... What tool do you use for capture ... I want to have an overview of the traffic ...) where NetFlow is the answer - or at least part of it
- Infrastructure is a potential target
- What's needed is actual insight

NetFlow flavors

- Several versions of the export protocol
- Metering and export keep being worked on
- Depends on vendor, hardware, software, and combinations of those
- Even when and where supported, particularities to keep in mind

Versions

- v5 - De facto standard (supported by non-C vendors implementing NetFlow)
- v9 - template-based
 - Flexibility (for now means that the user has to be very flexible) - complexity
 - May have some compelling features - SCTP, IPv6, some L2 fields, additional L3 (ttl, ipid), BGP next-hop, ...
 - But probably not processed by your collector of choice anyway
 - (Fast) moving target
 - Some features are backported to v5 (e.g. transport independency)
- v10 aka IPFIX - IETF-blessed standard (in its way to, at least)
 - Based on NetFlow v9 ("forked" from)
- Other - v1 (obsolete), v7 (switches), v8 (router-based aggregation)

Hardware peculiarities

- Support is not equal on all devices
 - No support on the switches below the cat 45xx
- Peculiarities where supported
 - On L3 switches, no TCP Flags, many specificities depending upon type of sup engine, PFC, config, software, ...

Hardware peculiarities - on routers too

Big **CAUTION** on every Cisco doco concerning NetFlow on the 12k

Entering this command on a Cisco 12000 Series Internet Router causes packet forwarding to stop for a few seconds while NetFlow reloads the route processor and line card CEF tables. To avoid interruption of service to a live network, apply this command during a change window, or include it in the startup-config file to be executed during a router reboot

Less advertised caution

Despite ASIC support in Engine 2, 3 and 4+ Linecards 'Full NetFlow' still inflicts a heavy burden on memory and therefore sampled netFlow is preferred

Outline

- 1 An introductory view of NetFlow
- 2 The flow cache
 - Overview
 - Populating the cache
 - Keeping the cache under control
 - Querying the cache in real time
 - Exporting the content of the cache
- 3 Collect & analysis
- 4 What's next

Principle

- A global memory cache, enabled for each interface
- number of entries - where it is kept
- For each packet entering an interface
 - Either an entry matching the key fields (saddr, daddr, sport, dport, L3 proto, ToS, input ifIndex) exists
 - Byte and packet counts are summed, end time is updated, TCP flags are ORed*
 - Or a new entry is created
- Expiration mechanism
- Export mechanism if so configured

Expiration mechanisms

When one of the following conditions is met

- Upon reception of a TCP *RST* or *FIN* packet
- After 15 seconds (default) for inactive flows
- After 30 minutes (default) for an active flow
- When the cache is full

Cisco, v5

Example

```
router(config)#ip flow-export version 5 [origin-as|peer-as]  
router(config-if)#ip flow {ingress | egress}
```

v9 specifics

Example

```
#ip flow-export version 9 [origin-as|peer-as] bgp-nexthop
#ip flow-capture ip-id
#ip flow-capture mac-addresses
#ip flow-capture packet-length
#ip flow-capture ttl
#ip flow-capture vlan-id
```

NDE - NetFlow Data Export

- NDE on the MSFC exports statistics for flows routed in software
- NDE on the PFC exports statistics for flows routed in hardware
- Flow masks on the PFC - from *source-only* to *full-interface*

Example

```
Switch(config)#mls netflow
Switch(config)#mls flow ip full
! To report L2 bridged traffic - not all PFCs
Switch(config)#ip flow export layer2-switched
Switch(config)#mls nde sender version 5
Switch(config-if)#ip flow {ingress | egress}
Switch(config)#ip flow-export destination 192.0.2.1 2055
```

cflowd on Juniper

- Relies on Service Cards / Adaptive Service Interfaces
- Configuration of a sampling rule in the firewall statement
- Rate of said sampling is configured in the forwarding-options
 - Can be 1
- Input filter calling the firewall rule as defined above for each interface

cflood on Juniper

Example

```
interfaces {
  sp-0/0/0 unit 0 family inet;
  fe-0/0/0 unit 0 family inet {
    filter input catch_all;
    address 10.88.17.126/28; }}
firewall {
  family inet filter catch_all term default then {
    sample; accept; }}
forwarding-options {
  sampling {
    input family inet {
      rate 1; max-packets-per-second 5000; }
    output { }}
```

FreeBSD

In-kernel ng_netflow(4) Netgraph node coupled with the ng_ipfw(4) ipfw Netgraph hook

Example

```
# ipfw add ngtee 10 ip from any to any
# ngctl mkpeer ipfw: netflow 10 iface0
# ngctl name ipfw:10 catchall
# ngctl msg catchall: setdlt { iface=0 dlt=12 }
# ngctl msg countall: settimeouts { inactive=3 active=300 }
```

FreeBSD cont'd.

Makes it easy to define flow export policies

Example

```
# ipfw add ngtee 5 tcp from any to me 22 in
# ngctl mkpeer ipfw: netflow 5 iface0
# ngctl name ipfw:5 ssh_in
# ngctl msg ssh_in: setdlt { iface=0 dlt=12 }
# ngctl msg ssh_in: settimeouts { inactive=3 active=300 }
```

- OOooohhHH
- Bad news is - AS fields are not populated - AAaahh...

- Tuning its size (32 to 256k depending on hw)
- Aging of entries
- Aggregation router side
- Sampling
- Input filtering
- Flow masks on Cat6.5k

Some performance considerations

- 64 bytes / entry in a cache
- Aggregation is more costly CPU-wise
- Done in hw for some platforms (Cat6.5k) - only export consumes CPU cycles
- Not fitted for more-than-reasonably busy routers
- Plan / test / monitor / prepare to rollback

Sampling or not

Depending on use / sf / hw

- Flow vs packet vs time - deterministic vs random
- Choice is not yours - check with your rep
- Effective (random sampled netflow) and scalable
- Granularity loss may not be an acceptable tradeoff for security
- Performance loss may not be an acceptable tradeoff for operations :D

```
#show ip cache flow
```

<i>SrcIf</i>	<i>SrcIPaddress</i>	<i>DstIf</i>	<i>DstIPaddress</i>	<i>Pr</i>	<i>SrcP</i>	<i>DstP</i>	<i>Pkts</i>
Vl90	10.11.10.50	Local	192.168.11.2	06	090D	0016	6
Vl97	192.168.19.214	Null	192.168.19.255	11	008A	008A	1
Vl16	192.168.17.165	Null	255.255.255.255	11	4001	006F	1
Vl10	192.168.15.65	Null	192.168.14.1	11	007B	007B	1
Vl10	192.168.15.166	Null	192.168.15.255	11	008A	008A	1
Vl10	192.168.15.84	Null	192.168.15.255	11	008A	008A	1
Vl10	192.168.15.75	Null	192.168.15.255	11	008A	008A	1

The top-talkers cache

- Possibility to create a dedicated cache for top-talkers
- Matching criteria, sorting order and number of entries are fixed at configuration
- Not the same as the *Dynamic Top Talkers* feature on the main cache

Example

```
router(config)#ip flow-top-talkers
router(config-flow-top-talkers)#top 30
router(config-flow-top-talkers)#sort-by bytes
```

The top-talkers cache, cont'd.

```
#show ip flow top-talkers
```

<i>SrcIf</i>	<i>SrcIPaddress</i>	<i>DstIf</i>	<i>DstIPaddress</i>	<i>Pr</i>	<i>SrcP</i>	<i>DstP</i>	<i>Bytes</i>
Se1/6:0	10.18.16.77	Fa0/1	10.11.19.88	06	05FD	0455	34K
Se1/6:0	10.18.16.77	Fa0/1	10.11.19.65	06	05FD	045D	10K
Se1/6:0	10.18.16.77	Fa0/1	10.11.14.51	06	05FD	04BA	9508
Se1/6:0	10.18.16.77	Fa0/1	192.168.9.7	06	05FD	059E	7064
Fa0/1	10.11.19.88	Se1/6:0	10.18.16.77	06	0455	05FD	5380

The JunOS way

```
#show services accounting flow-detail | match bgp | trim 26  
AAA.88.137.113 2146 Unknown BBB.88.137.126 bgp(179) 628 135136  
00:30:00 585 95087
```

FreeBSD

```
#flowctl catchall show
```

<i>SrcIf</i>	<i>SrcIPaddress</i>	<i>DstIf</i>	<i>DstIPaddress</i>	<i>Pr</i>	<i>SrcP</i>	<i>DstP</i>	<i>Pkts</i>
em0	10.19.11.2	lo0	10.18.5.89	1	0000	0000	51
(null)	10.18.5.89	em0	10.19.11.2	1	0000	0000	51
em0	10.21.5.19	lo0	10.18.5.89	1	0000	0000	8
(null)	10.18.5.89	em0	10.64.60.7	6	1466	9f8a	1

And the usual Unix tools | sort +n | head -n 10

Export mechanism

When expired, flows are packed together and sent to one or more collectors

- UDP based
- A header (24 bytes)
 - Version, number of *PDUs*, sequence number, ...
- 1-30 flow records (48 bytes each)
 - 1464-bytes packets
- Loss detection
 - SEQ number
- No provision for retransmission

On Cisco

Example

```
router(config)#ip flow-export destination 192.0.2.1 2055  
router(config)#ip flow-export source loopback0
```

Exporting the content of the cache, Juniper

Example

```
forwarding-options {  
  sampling {  
    output {  
cflowd 192.168.25.1 {  
  port 2055; version 5; autonomous-system-type origin; }  
interface sp-0/0/0 { source-address 192.168.25.2; };
```

On FreeBSD

Example

```
# ngctl mkpeer catchall: ksocket export inet/dgram/udp
# ngctl msg catchall:export connect inet/192.0.2.1:2055
```

Others

Soft meters / exporters

- PCAP to NetFlow: fprobe, softflowd, nProbe, ...
- Other sources: fw logs (PF, netfilter)

Improving the reliability of export

- Multiple (2) destinations where supported
- Multicast-aware collectors
 - Export destination can be a multicast address
 - Collectors join the multicast group
 - Out-of-the-box redundancy, scalability, easy testing of other collectors, collector diversity, ...
- Consider using other transport protocols - SCTP with backup ?
 - The way of the future as they say, IPFIX implementations *MUST* support SCTP
 - Though for now the list of collectors supporting SCTP is ... small (read: I don't know one, except NTOP being not a collector either)

Security considerations

- Unidirectional from *router to collector*
- UDP
- No crypto checksum
- No authentication of the exporter
- 48 bytes-long records for min 28-bytes UDP packets

You may want to protect the link between exporter and collector from a non invited participant

- ACLs, uRPF, TCP-Wrappers or alike

Outline

- 1 An introductory view of NetFlow
- 2 The flow cache
- 3 Collect & analysis**
 - Overview
 - Storage
 - Tools
 - nfdump / NfSen
 - direction & biflows
 - From PCAP to NetFlow
 - NetFlow and BGP
 - Other sources
 - Reductio ab aggregatio

Glories and duties of a collector

- Takes flows as input
 - May report loss
- Pre-processes
 - Deduplicates, long flows, filters, various aggregation schemes
- Stores
- Post-processes
 - Deduplicates, long flows, filters, various aggregation schemes
- Retrieval
 - Sorts, filters, aggregates,
- Graphs

Glories and duties of a collector

- Per record granularity
- Aggregation facilities and higher level views (while retaining granularity)
- Versatility
- Powerful CLI
- **Graphs**
- Can be splitted over different tools

Glories and duties of the analyst

- Look at graphs
- Look at TopNs
- Build filters to look at subsets of traffic
 - Control & management plane
 - Traffic targetting addresses in infrastructure subnets
 - Outgoing traffic from said addresses

More considerations

- IP-centric
 - Changing with v9: some L2, MPLS, ...
 - Not any-over-any friendly
- TCP flags are blended together
 - 'Isolated' flags are still worth looking at:
 - `nfdump -R . -o long 'flags S and not flags ARPFU'`
 - `nfdump -R . -o long 'flags R and not flags SPFU'`
- L7 streams are splitted upon 2+ records
- Source and destination does not match sockets' notion of client/server

Lost in collection

- Some figures: 3 STM1s, 2 borders, 1 (/1) sampling: 7 GB/day
- Mid term strategy: compression
- Long term: prune data and keep heavily aggregated summaries (topNs, ...)

Flat file vs database

- Db solution sounds sexy but is costly - per-record overhead
- B plan - keep granularity in "efficient" storages, and send aggregated and associated data in a db
- Probably no integration with used tools
- No normalized on-disk format

An overview of popular collectors

A lot in both the opensource and commercial world - varying a lot in scope and qualities (scoop)

- Arbor Peakflow SP, Lancope, Cisco MARS, Q1labs' QRadar, many more
- The venerable flow-tools, the yet more venerable CAIDA's cflowd, argus, SiLK, Stager, nfdump/NfSen, many more
- Others
 - aguri, glflow, panoptis, NTOP, stager, ...

Flow-tools / Flowscan

- Second historical collection-and-analysis package
- Many patches / associated scripts / ... / floating around
 - flow-tool to db, Flow Extract, ...
- Big user community
- Somewhat baroque (my opinion) CLI
- GUI - flowscan / cufLOW / jkflow
- Not actively maintained (latest rel 0505) - no v9 support

nfdump/NfSen

- Clean modular design
- Powerful and efficient (easy) yet simple (way harder) CLI with PCAP-like syntax
- tightly coupled GUI with nice graphs
- Takes input from NetFlow v5, v7, v9 (not all fields though) and sFlow
- Actively maintained
- Some plugins / patches floating around (porttracker, nfsplit, Holt-Winter)
- Community is growing steadily
- Versatile - troubleshooting, perf, security, ... - real time, trending, long-term analyses & forensics
- Comes with a flow-tools 2 nfdump converter
 - All your flows are belong to us

And the winner is ...

Architecture

nfcapd / sfcapd

- NetFlow / sFlow capture daemon
- No pre-post processing of the records collected
- Reads data from the network and writes to disk
 - Fixed-time binary files (default 5 mn)
- May demux to other collector

nfdump

- Reads files and displays records after run-time processing

Architecture

Others - self-explicit

- nfreplay, nfexpire, nfprofile, ft2nfdump

NfSen

- Graphs based on rrdtools
- Front-end to nfdump queries
- Plugins

Installation

Straightforward

- <http://sourceforge.net/projects/{nfdump,nfsen}>
- Stable branch vs snapshots
 - En-route to NfSen 1.3 - major usability improvements
 - Whatever the version, keep nfdump and nfsen synchronized
- Upgrade path provided between releases/snapshots
- Nfdump dependencies: c compiler
- Nfsen dependencies: perl, php, rrdtools, php-extensions (SESSION, SOCKETS)
- Bumping of max. number of SVIPC semaphores may be needed (one per collector)

nfdump

- Reads nfcapd files
- Processes
 - Parses pcap-like filters (usuals + if, packets, bytes, pps, bps, bpp, as, duration)
 - Matches long flows
 - Aggregates (proto, prefix/length, ASN, port)
 - TopNs (record, ip, port, tos, as, if, proto)
 - Orders by flow, packet, bytes, pps, bps, bpp
 - Anonymizes
- Displays
 - Pre-defined formats / user-defined formats
 - Statistics (number of flows, per proto, packets, bytes, ...)

nfdump output

- Pre-defined formats: `-o "line|long|extended|pipe|raw|fmt"`
- `-o "fmt:%ts %td %sap %dap"`

```
-r nfcapd.200701202315 -c 3 -o "fmt:%td %sap %dap"
```

Duration	Src IP Addr:Port	Dst IP Addr:Port
0.000	BBB.158.102.14:53	AAA.189.5.89:53
11.000	CCC.249.66.142:52812	AAA.189.5.89:80
1.000	DDD.218.231.90:4509	AAA.189.5.90:445

AS Matrix

```
nfdump -M <source_list> \  
-R nfcapd.$start_tslot:nfcapd.$end_tslot \  
-s record/bytes -A srcas,dstas -n 0 \  
-o "fmt:%sas %das %byt"
```

If user-defined is not enough - nfdump.c

Define your output format string.

Test the format using standard syntax `-o "fmt:<your format>"`

Create a `#define` statement for your output format, similar than the standard output formats above.

Add another line into the `printmap[]` struct below

BEFORE the last NULL line for you format:

```
{ "formatname", format_special, FORMAT_definition, NULL },
```

The first parameter is the name of your format as recognized on the command line as `-o <formatname>`

The second parameter is always `'format_special'` - the printing function.

The third parameter is your format definition as defined in `#define`.

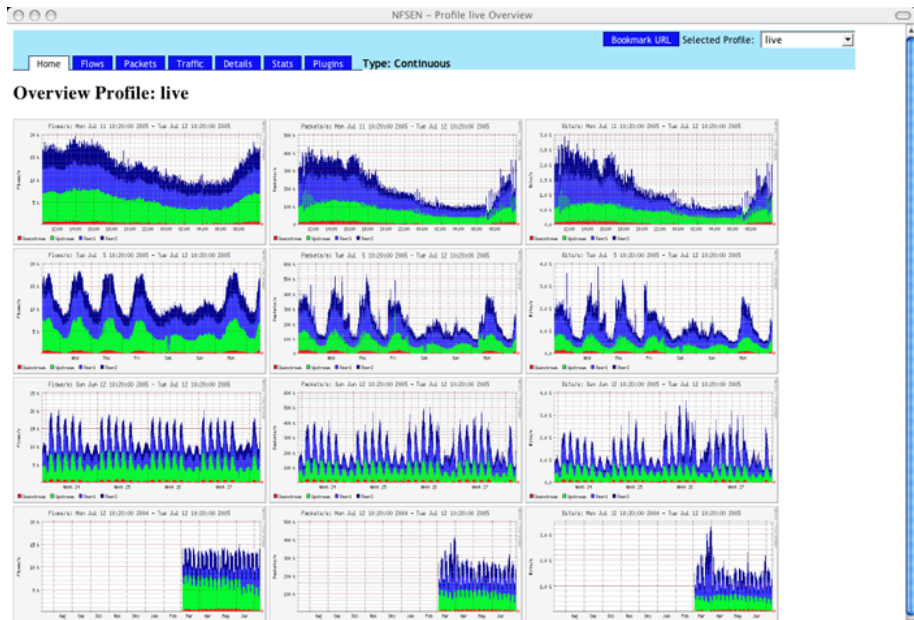
The forth parameter is always NULL for user defined formats.

Recompile nfdump

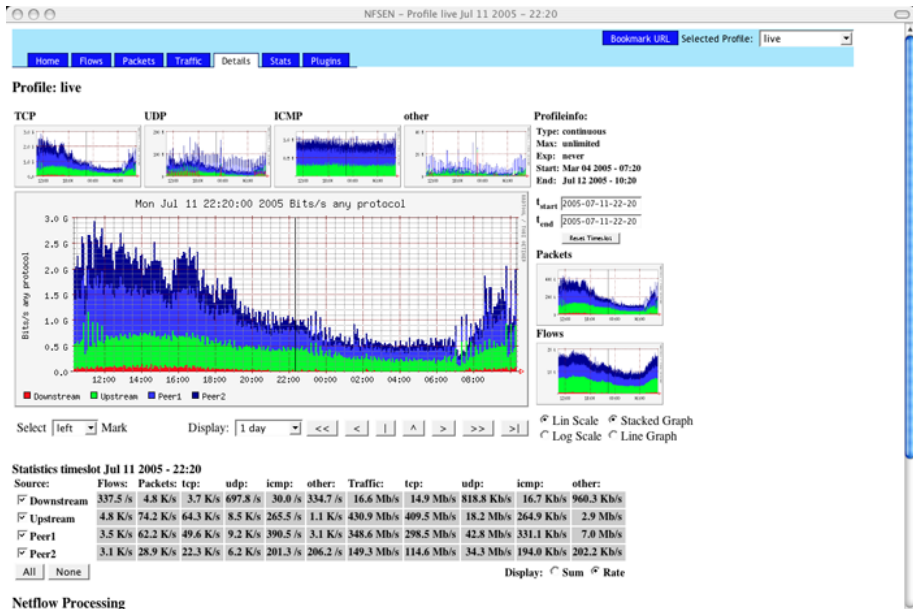
\$CONFDIR/nfsen.conf

- Various paths (\$BASEDIR, \$BINDIR, \$LIBEXECDIR, \$CONFDIR, \$HTMLEDIR, \$VARDIR, \$PROFILESTATDIR, \$PROFILEDATADIR, \$BACKEND_PLUGINDIR, \$FRONTEND_PLUGINDIR, \$PREFIX
- User of the {s,n}fcpd processes
- Data layout (see below)
- Disk filling % (warning message)
- Disk filling water marks for profiles
- Sources
 - \$BINDIR/nfsen reconfig to add more

Main tab



Navigation tab



Data organization

- Sources - aka exporters
- nfsplit
 - External contrib
 - Stands between the actual flows and nfcapd
 - Splits per interface
- Data lives in \$BASEDIR/profiles/<profile>/<source>

```
# 0 default      no hierachy levels - flat layout - compatible
# 1 %Y/%m/%d     year/month/day
# 2 %Y/%m/%d/%H  year/month/day/hour
# 3 %Y/%W/%u     year/week_of_year/day_of_week
# 4 %Y/%W/%u/%H  year/week_of_year/day_of_week/hour
# 5 %Y/%j        year/day-of-year
# 6 %Y/%j/%H     year/day-of-year/hour
# 7 %Y-%m-%d     year-month-day
# 8 %Y-%m-%d/%H  year-month-day/hour
```

Data organization, cont'd

- Profiles - 'live' post-install
 - Managed through web *and* CLI
 - Data is duplicated between profiles (should change)
 - Can be retroactive
 - Time-delimited or continuous
 - Expiration of old data - based on duration and/on size
 - Beware - thresholds are per profile, no global check
 - Beware - it is recommended to put the stats (\$PROFILESTATDIR) and the data (\$PROFILEDATADIR) on different volumes to prevent corruption in case of a disk full
- Channels - defined by filters

Managing profiles

- List all profiles

```
$BINDIR/nfsen -A
```

- Add a profile

```
$BINDIR/nfsen -a <profile> -c desc -B <starts> -E <ends>
```

- Modify a profile

```
$BINDIR/nfsen -m <profile>
```

- Delete a profile

```
$BINDIR/nfsen -d <profile>
```

Managing profiles

```
/data/nfsen/bin/nfsen -a slammer -B 2006-10-12-23-45 -S  
other_in:other_out 'proto udp and port 1434'
```

```
#  
name      slammer  
tstart    Thu Oct 12 23:45:00 2006  
tend      Wed Dec 6 02:55:00 2006  
updated   Thu Oct 12 23:40:00 2006  
filter    filter.txt  
expire    0 hours  
size      0  
maxsize   0  
sources   other_in:other_out  
type      continuous  
locked    1  
status    new
```

Managing profiles

```
/data/nfsen/bin/nfsen -l slammer
```

```
#  
name      slammer  
tstart    Thu Oct 12 23:45:00 2006  
tend      Wed Dec 6 02:55:00 2006  
updated   Thu Oct 12 23:40:00 2006  
filter    filter.txt  
expire    0 hours  
size      0  
maxsize   0  
sources   other_in:other_out  
type      continuous  
locked    1  
status    built 53.9%
```

Navigation tab


Home | **Flows** | Packets | Traffic | Details | Stats | Plugins | Type: Continuous

Bookmark URL Selected Profile: slammer

Overview Profile: slammer



Channels



Detail tab

Netflow Processing

List Flows Stat TopN

Source: Filter:

port 22

Other
win
icecast
ssh
dns_out
dns_in

All Sources and <none>

Options:

Top: 10

Stat: Flow Records order by flows

Aggregate ☒ proto ☒ srcPort srcIP
☒ dstPort dstIP

Limit: ☐ Packets > 0 -

Output: line ☐ / IPv6 long

```
nfdump -M /data/nfsen/profiles/./main/Other:win:icecast:ssh:dns_out:dns_in:smtp_out:smtp_in:http_out:http_in:jabber:bgp -R nfcapd.200702021945:nfcapd.200702022345 -n 10 -s
```

Aggregated flows 4545

Top 10 flows ordered by flows:

Date flow start	Duration	Proto	Src IP Addr:Port	Dst IP Addr:Port	Packets	Bytes	Flows
2007-02-02 21:41:49.132	6380.976	TCP	220.228.154.118:36862	220.228.154.118:51813			6
2007-02-02 20:38:07.753	9330.649	TCP	220.228.154.118:51813	220.228.154.118:34672			6
2007-02-02 21:24:05.863	5695.911	TCP	220.228.154.118:34672	220.228.154.118:50383			6
2007-02-02 20:04:03.523	1513.209	TCP	220.228.154.118:50383	220.228.154.118:41773			6
2007-02-02 21:09:05.227	6731.942	TCP	220.228.154.118:41773	220.228.154.118:45885			6
2007-02-02 21:18:54.574	5686.970	TCP	220.228.154.118:45885	220.228.154.118:44939			6
2007-02-02 20:10:48.185	10830.607	TCP	220.228.154.118:44939	220.228.154.118:46214			6
2007-02-02 20:44:56.447	7200.245	TCP	220.228.154.118:46214	220.228.154.118:41513			6
2007-02-02 21:00:30.717	7781.812	TCP	220.228.154.118:41513	220.228.154.118:51682			6
2007-02-02 19:56:08.329	11295.841	TCP	220.228.154.118:51682				6

220.228.154.118:
LL-220-228-154-118.LL.sparqnet.net

AD | IP | AD Name
9919 | 220.228.154.118 | BCC-TV New Century InfoComm

Summary: total flows: 9755, total bytes: 6.1 M, total packets: 62688, avg bps: 3472, avg pps: 4, avg bpp: 102
Time window: 2007-02-02 19:33:24 - 2007-02-02 23:49:54
Total flows processed: 22738, skipped: 0, Bytes read: 1188184
Sys: 1.299s flows/second: 17580.9 Wall: 0.048s flows/second: 472134.6

AS resolution

```
$BASEDIR/libexec/Lookup.pm
```

```
my $whois_socket = IO::Socket::INET->new(  
    PeerAddr  => 'whois.cyberabuse.org',  
    PeerPort  => 43,  
    Proto     => 'tcp',  
    timeout   => 10 );
```

Plugins

- NfSen is "plugin" friendly
- Backend in perl and frontend in PHP
- Skeleton available in \$BASEDIR/plugins/demoplugin.pm and \$HTMLDIR/plugins/demoplugin.php
- No repository of plugins
 - Only publicly available plugin is PortTracker
 - Tracks ... TopN ports

Holt-Winter patches


- hw-patched rrdtools for NfSen
- External contrib
- Discussions pending regarding integration or not
- Can be installed in parallel to an existing NfSen installation
- Useful to spot traffic irregularities, especially among more verbose sources

Holt-Winter - main view

Home Flows Packets Traffic Details Stats Plugins


Bookmark URL Selected Profile: live

Profile: live



Profileinfo:

Type: continuous
Max: unlimited
Exp: never
Start: Aug 11 2006 - 11:55
End: Oct 13 2006 - 09:05




tstart: 2006-10-12-21-05
tend: 2006-10-12-21-05
[Reset Timeslot](#)

Packets



Traffic



Anomaly view

[Home](#)
[Flows](#)
[Packets](#)
[Traffic](#)
[Details](#)
[Stats](#)
[Plugins](#)

Bookmark URL

Selected Profile: live




Profile: live

TCP




UDP



ICMP



other



Profileinfo:


Type: continuous

Max: unlimited


Exp: never

Start: Aug 11 2006 - 11:55



End: Oct 13 2006 - 09:10



Packets



Traffic



The example below is purely fictional

Example

```
+-----+
| 1. The packet contains a specific |
| crafted IP option.                |
|-----|
| AND                               |
|-----|
| 2. The packet is one of the following |
| protocols:                         |
|-----|
| * ICMP - Echo (Type 8) - 'ping'    |
|-----|
| * ICMP - Timestamp (Type 13)       |
|-----|
| * ICMP - Information Request (Type |
| 15)                                |
|-----|
```


Example

```
+-----+
|  * ICMP - Address Mask Request (Type |
|  17)                                |
|-----|
|  * PIMv2 - IP protocol 103          |
|-----|
|  * PGM - IP protocol 113            |
|-----|
|  * URD - TCP Port 465                |
|-----|
|  AND                                |
|-----|
|  3. The packet is sent to a physical |
|  or virtual IPv4 address configured on |
|  the affected device.                 |
+-----+
```

The ICMP case

- Source port == 0
- Code+type shoehorned into dst port - code is lower 8 bits, type higher 8 bits

tstamp request & tstamp reply

ICMP	192.168.2.11:0	->	AAA.189.5.89:13.0
ICMP	AAA.189.5.89:0	->	192.168.2.11:14.0

- No specific filters for ICMP though

Type 8, 13, 15, 17

proto icmp and (port 2048 or port 3328 or port 3840 or port 4352)

Adding a new profile and channel

Home Graphs Details Alerts Stats Plugins Continuous Profile [Bookmark URL](#) Profile: **blah** ▼

Building Profile: blah

7.6%

Group: (nogroup)

Description: blah

Type: Continous

Start: 2007-01-12-14-45

End: 2007-02-04-11-40

Last Update: 2007-01-12-14-40

Size: 0 B

Max. Size: unlimited

Expire: never

Status: built 7.6% - locked

▼ Channellist:

▼ adv2

Colour: #FF66FF Sign: + Order: 1

Filter: (proto icmp and (port 2048 or port 3328 or port 3840 or port 4352)) or proto 103 or proto 113 or (proto tcp and port 465)

Sources: source1

Analyst's temptation

- Looking at 'streams' rather than flows

Example

```
stime, client:sport, sflags, sbytes, spkts -> server:dport,  
dflags, dbytes, dpkts, ...
```

- Makes it easier to track some irregularities / asymmetries
 - SYN / RST couples
 - Unusual proportions of s{bytes, pkts, ...} vs d{bytes, pkts, ...}
- Easier classification of L7 when knowing connect() and bind() sides
- Orthogonal problems

Direction

- Records convey no direction information
- No 3-way handshake granularity with TCP flags
- No sub-ms granularity for start times
- Besides, many reasons to have timestamps in the 'wrong' order
- No good heuristics beyond
 - Parse time-sorted
 - Swap saddr and daddr if
$$\text{sport} < 1025 \ \&\& \ \text{dport} > 1024$$
- Still too many incorrectly headed flows

Biflows

- Basically 3 places to aggregate 2 flows into 1 biflow
 - Router-side
 - Collect - one shot - increases complexity
 - post-process - run-time
- Some IPFIX drafts

Who'd want that ?

- Everybody trying to make sense of a (big, for some meaning of) packet capture
- In need of efficient tools
- Re-use of existing tools / methods

Argus

- <http://qosient.com/argus/>
- Collect and analysis of network data
 - PCAP (live or off-line)
 - NetFlow
- Stores flow-like data
 - Bidirectionnal records


The other way around ? (from NetFlow to PCAP)

- Nobody would want that
- Lot of tools working on PCAP data where NetFlow records would be fine
 - Specialized reports, DoS, ...
- No normalized on-disk format - no incentive for tools' authors
- Disclaimer: very clumsy and inelegant
- `nfdump` -> ASCII -> `ipsumdump` -> PCAP
 - OUCH
- <http://www.cs.ucla.edu/~kohler/ipsumdump/>

BGP

- Heartbeat-like traffic
 - Whose level of steroids in blood increases over time, but this is another story
- Pinpointing of BGP events (which scale ?)
- At least timeframe is provided - up to the ops to dig into BGP logs

BGP heartbeat




eBGP multihop instability

Flows

Packets

Traffic

Profile: bgp, Group: (nogroup) - traffic




eBGP multihop instability

Flows

Packets



Traffic

Profile: bgp, Group: (nogroup) - packets



Looking at the noise

- Sink Holes
- Non used address space
- Diverted traffic



Aguri

- <http://www.csl.sony.co.jp/kjc/software.html#aguri>
- Automatic traffic aggregation - prefixes & ports
- PCAP-based

Aguri

aguri

```
[src address] 1049390337 (100.00%)
                AAA.64.60.71      1008893224 (96.14%)
                BBB.189.5.89      36286986 (3.46%)
[dst address] 1049390337 (100.00%)
                AAA.64.60.71      35493322 (3.38%)
                BBB.189.5.89      1012495103 (96.48%)
[ip:proto:srcport] 1049390337 (100.00%)
                4:6:22  1008772431 (96.13%)
                4:6:62008      35303106 (3.36%)
[ip:proto:dstport] 1049390337 (100.00%)
                4:6:22  35439904 (3.38%)
                4:6:62008      1008498420 (96.10%)
```

Outline

- 1 An introductory view of NetFlow
- 2 The flow cache
- 3 Collect & analysis
- 4 What's next
 - Perhaps not really next
 - Resources

NetFlow v9 and v10

- Disclaimer: I'm not affiliated with ...
- Disclaimer: check with your reps
- Templates ...
- Other transports
- User-defined caches and export policies
- Other vendors should follow

Dissemination of (net)flow specification rules

- Disclaimer
- draft-marques-idr-flow-spec-03
- "Successor" of uRPF + BGP null routes
- Former is a clever engineering trick to null route prefixes
- Latter formalizes the propagation of flow-like informations through BGP for further action
- Dst prefix, src prefix, proto, src & dst port, icmp type & code, TCP flags, packet length, ToS, fragment
- Basically a flow record

State of the flow-spec draft

- Implemented by one vendor
- Other is waiting for customer traction ?
- A NetFlow vendor from MI
- No (public) implementation on any Unix BGP daemon - too bad
 - Spot traffic to be acted upon from a BGP-speaking NetFlow analysis workstation
 - Use flow-spec support to inject and propagate said records

Some links

- <http://www.cisco.com/go/netflow/>
- http://www.cisco.com/en/US/products/ps6601/prod_presentation_list.html
- <http://www.switch.ch/tf-tant/floma/software.html>